

Custom PII Detection Guardrail Implementation Guide

Step-by-step guide for implementing my custom guardrail on litellm and do local development for image building to push to ECR

Technical Assessment Requirements

This guide addresses the specific technical objectives:

1. *Deploy LiteLLM*: with at least one AI model (local container deployment)
2. *Develop custom guardrails*: to detect and block PII (Email + SSN + Credit Card + Phone Number)
3. *Configure LiteLLM*: to use custom guardrails (Regex and Microsoft Presidio)
4. *Understanding Architecture*: with architecture diagram
5. *Demonstrate functionality*: By showing the integrated guardrail into LiteLLM

Prerequisites

- Docker and Docker Compose installed (local development purposes)
- Git installed (version control)
- Basic command line familiarity (enough to work with docker)
- Text editor or IDE (local dev environment for image development)

Step 1: Clone and Setup the Repository

```
```bash
Clone the repository
git clone https://github.com/mrcloudchase/litellm-app.git
cd litellm-app

Verify the structure
ls -la
```
```

Expected files:

- `Dockerfile` - Container build instructions
- `docker-compose.yml` - Local development stack
- `litellm-config-local.yaml` - LiteLLM configuration used by Docker Compose for local development
- `pii_*.py` - Custom guardrail implementations
- `Makefile` - Automation commands (automates docker compose up/down, etc.)

Step 2: Deploy LiteLLM with Local AI Model

Build and Start the Stack

```
```bash
Build all containers
make build

Start the complete stack (LiteLLM + Ollama + PostgreSQL)
make start

Pull a local AI model for testing llama3.2:3b
make pull-model

Verify everything is running
docker ps
```
```

Expected output:

```
```
CONTAINER ID IMAGE PORTS
abc123... litellm-app-litellm 0.0.0.0:4000->4000/tcp
def456... litellm-app-ollama 0.0.0.0:11434->11434/tcp
ghi789... postgres:15-apline 0.0.0.0:5432->5432/tcp
```
```

Verify Base Deployment

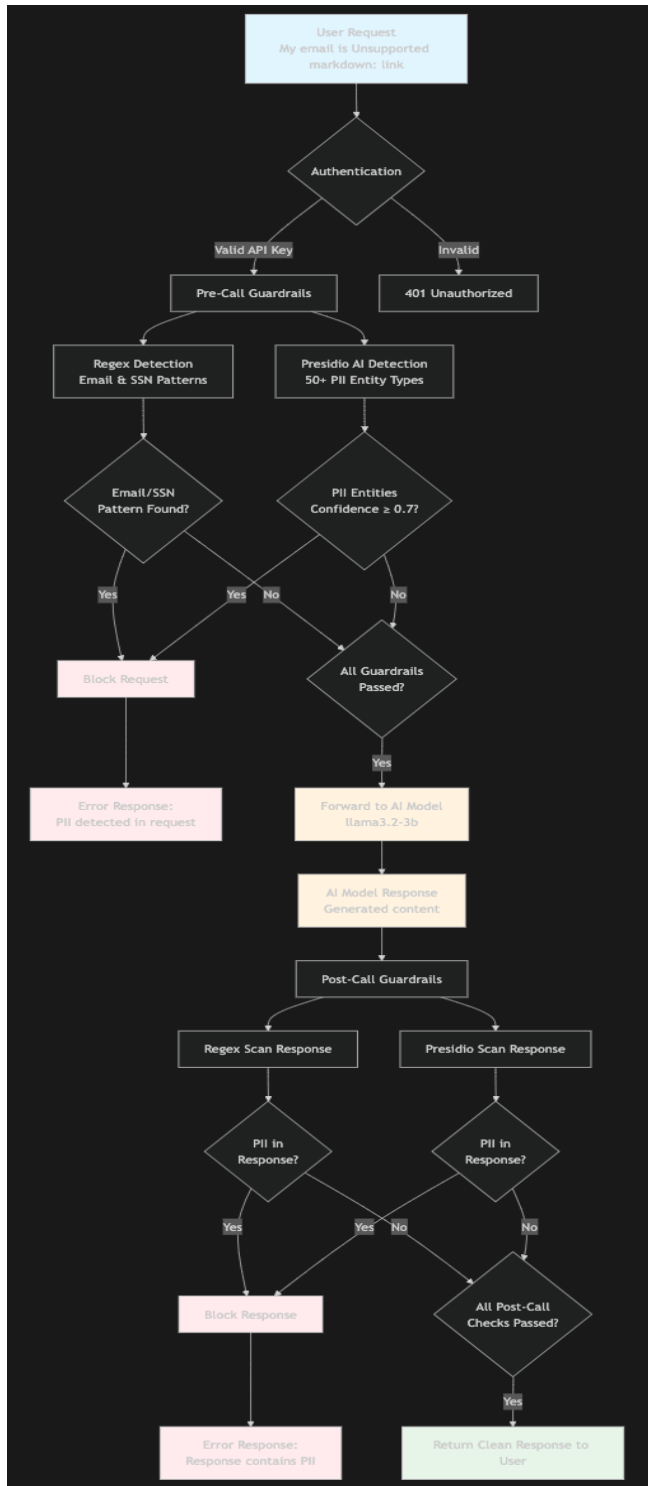
```
```bash
Test LiteLLM is responding
curl -X GET http://localhost:4000/health

Test AI model integration (without guardrails)
curl -X POST http://localhost:4000/v1/chat/completions \
 -H "Authorization: Bearer sk-local-dev-key-12345" \
 -H "Content-Type: application/json" \
 -d '{
 "model": "llama3.2-3b",
 "messages": [{"role": "user", "content": "Hello, how are you?"}],
 "max_tokens": 10
 }'
```
```

Expected response: JSON with AI-generated content

Step 3: Understanding the Custom Guardrail Implementation

Architecture Overview



Key Implementation Files

1. ``pii_regex_detection.py`` - Shared regex patterns for email and SSN detection
2. ``pii_regex_precall.py`` - Pre-call guardrail (blocks PII in user input)
3. ``pii_regex_postcall.py`` - Post-call guardrail (blocks PII in AI responses)
4. `**`pii_presidio*.py`**` - AI-powered guardrails using Microsoft Presidio
5. ``litellm-config.yaml`` - Configuration that registers the guardrails

Core Guardrail Logic

Email Detection Pattern:

```
```python
email_pattern = re.compile(r'\b[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}\b')
```

### **SSN Detection Pattern:**

```
```python
ssn_pattern = re.compile(r'\b\d{3}[-\s]?\d{2}[-\s]?\d{4}\b')
```

Step 4: Test the Custom Guardrails

Test 1: Baseline (No Guardrails)

```
```bash
Test without any guardrails - should work normally
curl -X POST http://localhost:4000/v1/chat/completions \
 -H "Authorization: Bearer sk-local-dev-key-12345" \
 -H "Content-Type: application/json" \
 -d '{
 "model": "llama3.2-3b",
 "messages": [{"role": "user", "content": "My email is
test@example.com"}],
 "max_tokens": 10
 }'
```
```

Expected: Normal AI response (guardrails not active)

Test 2: Email Detection with Regex Guardrail

```
```bash
Test email detection - should be BLOCKED
curl -X POST http://localhost:4000/v1/chat/completions \
 -H "Authorization: Bearer sk-local-dev-key-12345" \
 -H "Content-Type: application/json" \
 -d '{
 "model": "llama3.2-3b",
 "messages": [{"role": "user", "content": "My email is
test@example.com"}],
 "guardrails": ["pii-regex-precalls"],
 "max_tokens": 10
 }'
```
```

Expected response:

```
```json
{
 "error": {
 "message": "Pre-call guardrail blocked PII detected: email",
 "type": "BadRequestError",
 "code": 400
 }
}
```
```

Test 3: SSN Detection with Regex Guardrail

```
```bash
Test SSN detection - should be BLOCKED
curl -X POST http://localhost:4000/v1/chat/completions \
 -H "Authorization: Bearer sk-local-dev-key-12345" \
 -H "Content-Type: application/json" \
 -d '{
 "model": "llama3.2-3b",
 "messages": [{"role": "user", "content": "My SSN is 123-45-6789"}],
 "guardrails": ["pii-regex-precalls"],
 "max_tokens": 10
 }'
```
```

Expected response:

```
```json
{
 "error": {
 "message": "Pre-call guardrail blocked PII detected: ssn",
 "type": "BadRequestError",
 "code": 400
 }
}
```
```

Test 4: Clean Input (Should Pass)

```
```bash
Test clean input - should work normally
curl -X POST http://localhost:4000/v1/chat/completions \
 -H "Authorization: Bearer sk-local-dev-key-12345" \
 -H "Content-Type: application/json" \
 -d '{
 "model": "llama3.2-3b",
 "messages": [{"role": "user", "content": "Hello, how are you today?"}],
 "guardrails": ["pii-regex-precalls"],
 "max_tokens": 10
 }'
```
```

Expected: Normal AI response (no PII detected)

Test 5: AI-Powered Presidio Guardrail (Stretch Goal)

```
```bash
Test comprehensive PII detection with AI
curl -X POST http://localhost:4000/v1/chat/completions \
 -H "Authorization: Bearer sk-local-dev-key-12345" \
 -H "Content-Type: application/json" \
 -d '{
 "model": "llama3.2-3b",
 "messages": [{"role": "user", "content": "My name is John Smith and I
work at Microsoft in Seattle"}],
 "guardrails": ["pii-presidio-precalls"],
 "max_tokens": 10
 }'
```
```

Expected response:

```
```json
{
 "error": {
 "message": "Pre-call guardrail blocked PII detected: PERSON,
ORGANIZATION, LOCATION",
 "type": "BadRequestError",
 "code": 400
 }
}
```
```

Step 5: Use HTTP Test Collections (Alternative Testing Method)

Using VS Code REST Client

1. ****Install REST Client extension**** in VS Code
2. ****Open test files:****
 - `tests/test_regex.http` - Regex guardrail tests
 - `tests/test_presidio.http` - Presidio guardrail tests
3. ****Run tests interactively:****
 - Click "Send Request" above each test
 - Observe responses in VS Code

Using the Makefile

```
```bash
Run automated tests
make test-guardrails

View test results
cat test_results.log
```
```

Step 6: Examine the Implementation Code

Regex Detection Engine (`pii_regex_detection.py`)

```
```python
import re
from typing import Dict, List

class PIIRegexDetection:
 def __init__(self):
 self.patterns = {
 'email': re.compile(r'\b[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}\b'),
 'ssn': re.compile(r'\b\d{3}[-\s]?\d{2}[-\s]?\d{4}\b'),
 'phone': re.compile(r'\b(?:[0-9]{3}\)?[-.\s]?[0-9]{3}[-.\s]?[0-9]{4}\b'),
 'credit_card': re.compile(r'\b(?:\d{4}[-\s]?){3}\d{4}\b')
 }

 def detect_pii(self, text: str) -> Dict[str, List[str]]:
 """Detect PII in text and return findings"""
 detected = {}
 for pii_type, pattern in self.patterns.items():
 matches = pattern.findall(text)
 if matches:
 detected[pii_type] = matches
 return detected

 def has_pii(self, text: str) -> bool:
 """Quick check if text contains any PII"""
 return bool(self.detect_pii(text))
...
```
```


Pre-Call Guardrail (`pii_regex_precall.py`)

```
```python
from litellm.integrations.custom_guardrail import CustomGuardrail
from pii_regex_detection import PIIRegexDetection

class PIIRegexPreCallGuardrail(CustomGuardrail):
 def __init__(self, **kwargs):
 self.detector = PIIRegexDetection()
 super().__init__(**kwargs)

 async def async_pre_call_hook(self, user_api_key_dict, cache, data,
call_type):
 # Extract user messages
 user_content = ""
 for message in data.get("messages", []):
 if message.get("role") == "user":
 user_content += message.get("content", "") + " "

 # Detect PII
 detected_pii = self.detector.detect_pii(user_content)

 if detected_pii:
 pii_types = list(detected_pii.keys())
 raise Exception(f"Pre-call guardrail blocked PII detected: {'',
'.join(pii_types)}")

 return None # Allow request to proceed
```
```

Configuration Registration (`litellm-config.yaml`)

```
```yaml
guardrails:
 # Regex-based guardrails (fast, pattern-matching)
 - guardrail_name: "pii-regex-precalls"
 litellm_params:
 guardrail: pii_regex_precall.PIIRegexPreCallGuardrail
 mode: "pre_calls"

 - guardrail_name: "pii-regex-postcalls"
 litellm_params:
 guardrail: pii_regex_postcall.PIIRegexPostCallGuardrail
 mode: "post_calls"

 # AI-powered guardrails (comprehensive, context-aware)
 - guardrail_name: "pii-presidio-precalls"
 litellm_params:
 guardrail: pii_presidio_precall.PIIPresidioPreCallGuardrail
 mode: "pre_calls"
 language: "en"
 threshold: 0.7
```
```

Step 7: Demonstrate Advanced Features

Multiple Guardrails

```
```bash
Use both regex and AI guardrails together
curl -X POST http://localhost:4000/v1/chat/completions \
 -H "Authorization: Bearer sk-local-dev-key-12345" \
 -H "Content-Type: application/json" \
 -d '{
 "model": "llama3.2-3b",
 "messages": [{"role": "user", "content": "Contact me at john@company.com
or call 555-123-4567"}],
 "guardrails": ["pii-regex-precalls", "pii-presidio-precalls"],
 "max_tokens": 10
 }'
```
```

Post-Call Protection

```
```bash
Test post-call guardrail (if AI model tries to output PII)
curl -X POST http://localhost:4000/v1/chat/completions \
 -H "Authorization: Bearer sk-local-dev-key-12345" \
 -H "Content-Type: application/json" \
 -d '{
 "model": "llama3.2-3b",
 "messages": [{"role": "user", "content": "Generate a fake email address
for testing"}],
 "guardrails": ["pii-regex-postcalls"],
 "max_tokens": 50
 }'
```
```

Step 8: Customize for Your Needs

Adding New PII Types

1. Edit `pii_regex_detection.py`:

```
```python
self.patterns = {
 'email': re.compile(r'\b[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}\b'),
 'ssn': re.compile(r'\b\d{3}[-\s]?\d{2}[-\s]?\d{4}\b'),
 'phone': re.compile(r'\b(?:[0-9]{3}\)?[-.\s]?[0-9]{3}[-.\s]?[0-9]{4}\b'),
 'credit_card': re.compile(r'\b(?:\d{4}[-\s]?){3}\d{4}\b'),
 # Add your custom pattern here
 'ip_address': re.compile(r'\b(?:[0-9]{1,3}\.){3}[0-9]{1,3}\b'),
}
```
```

2. Rebuild and test:

```
```bash
make build
make start
Test your new pattern
```
```

Configuring Presidio Entities

1. Edit `pii_presidio_detection.py`:

```
```python
self.entities = [
 "PERSON", "EMAIL_ADDRESS", "PHONE_NUMBER", "ORGANIZATION",
 "LOCATION", "IP_ADDRESS", "CREDIT_CARD", "US_SSN",
 # Add or remove entities as needed
 "US_PASSPORT", "MEDICAL_LICENSE", "CRYPTO"
]
```
```

Step 9: Production Deployment Preparation

Container Registry Push

```
```bash
Build for production
docker build -t your-registry/litellm-guardrails:latest .

Push to your registry
docker push your-registry/litellm-guardrails:latest
```
```

Environment Configuration

```
```bash
Set production environment variables
export LITELLM_MASTER_KEY="your-secure-master-key"
export LITELLM_MODE="PRODUCTION"
export DATABASE_URL="postgres://user:pass@host:5432/litellm"

Deploy to your infrastructure
docker run -p 4000:4000 \
 -e LITELLM_MASTER_KEY=$LITELLM_MASTER_KEY \
 -e LITELLM_MODE=$LITELLM_MODE \
 -e DATABASE_URL=$DATABASE_URL \
 your-registry/litellm-guardrails:latest
```
```

Step 10: Walkthrough Preparation

Key Demonstration Points

1. Show the running deployment:

```
```bash
curl http://localhost:4000/health
docker ps
```
```

2. Demonstrate PII blocking:

- Email detection with regex guardrail
- SSN detection with regex guardrail
- Person/organization detection with Presidio

3. Show clean requests passing through:

- Normal conversation without PII
- AI model responding normally

4. Explain the code structure:

- Detection engine separation
- Pre-call vs post-call hooks
- Configuration in YAML

5. Discuss the architecture:

- Dual detection strategy (fast + comprehensive)
- Container-based deployment
- Production readiness features

Troubleshooting Common Issues

Issue 1: Container Won't Start

```
```bash
Check logs
docker logs litellm-app-litellm-1

Common fix: ensure ports aren't in use
sudo lsof -i :4000
sudo lsof -i :11434
```
```

Issue 2: Model Not Found

```
```bash
Ensure Ollama model is pulled
docker exec litellm-app-ollama-1 ollama list

Pull the model if missing
make pull-model
```
```

Issue 3: Guardrail Not Triggering

```
```bash
Check configuration
cat litellm-config.yaml

Verify guardrail is specified in request
"guardrails": ["pii-regex-precalls"]
```
```

Issue 4: Authentication Errors

```
```bash
Verify master key
export MASTER_KEY="sk-local-dev-key-12345"

Use in requests
curl -H "Authorization: Bearer $MASTER_KEY" ...
```
```

Success Criteria Checklist

Core Requirements Met:

- [] **LiteLLM deployed** with local container and AI model
- [] **Custom guardrail developed** with clean, readable code using classes/methods
- [] **Two PII types detected**: Email addresses and SSNs
- [] **Blocks on input and output**: Pre-call and post-call protection
- [] **LiteLLM configured** to use the custom guardrails
- [] **Documentation created** with architecture diagram
- [] **Functionality demonstrated** with working examples

Stretch Goal Met:

- [] **Microsoft Presidio integrated** for comprehensive PII detection
- [] **50+ entity types** supported beyond just email/SSN
- [] **Context-aware detection** using ML models
- [] **Configurable confidence thresholds** for AI detection

Professional Implementation:

- [] **Production-ready code** with error handling
- [] **Comprehensive testing** with HTTP test collections
- [] **Container deployment** with health checks
- [] **Configuration management** via YAML
- [] **Performance optimization** with dual detection strategy
- [] **Documentation** suitable for technical audience

1. **Deployed LiteLLM** with Ollama AI model in containers
2. **Custom PII guardrails** detecting emails, SSNs, and 50+ other entity types
3. **Dual protection** system (regex + AI) for comprehensive coverage
4. **Production-ready** implementation with proper architecture
5. **Comprehensive documentation** with step-by-step instructions
6. **Working demonstration** ready for technical walkthrough

The implementation showcases **enterprise-grade PII protection** with clean, maintainable code that can be easily extended for additional use cases.